

# Discussion 6

# Introduction to Tensorflow.Keras

EE599 Deep Learning  
Kuan-Wen (James) Huang  
Spring 2020

# Installation

- Activate your virtual environment, either Anaconda or pyenv
- Install Tensorflow via
  - `pip install tensorflow` (works for both Anaconda or pyenv)
  - Install tensorflow package with Anaconda-navigator
- For GPU support (available for Ubuntu and Windows with CUDA<sup>®</sup>-enabled cards), check <https://www.tensorflow.org/install/gpu>

# Guides and References of tf.Keras

- Guide:  
<https://www.tensorflow.org/guide/keras/overview>
- Tutorial:  
<https://www.tensorflow.org/tutorials>
- Reference:  
[https://www.tensorflow.org/api\\_docs/python/tf/keras/](https://www.tensorflow.org/api_docs/python/tf/keras/)

# Two ways of using tf.Keras

- Sequential model
  - Easy
  - Less freedom
- Functional API model
  - Good for complex models
    - Models with branches, multiple inputs or multiple outputs
    - Layers with shared parameters
    - Custom layer, output
    - Custom functions

# Sequential Model – example 1

- Check [1\\_fmnist.py](#)
- Create a model with `keras.Sequential()` with a list of layer objects including input, hidden layers and output, e.g.

```
model = keras.Sequential([  
    keras.layers.Input(shape=(784,), name='input'),  
    keras.layers.Dense(128, activation='relu', name='hidden1'),  
    keras.layers.Dense(10, activation='softmax', name='output')  
])
```

# Sequential Model – example 2

```
model = keras.Sequential()  
model.add(keras.Layers.Dense(128, input_shape=(784,)))  
  
# now the model will take as input arrays of shape (*, 784)  
# and output arrays of shape (*, 128)  
# after the first layer, you don't need to specify  
# the size of the input anymore  
  
model.add(keras.Layers.Activation('relu'))  
# add an activation layer  
  
model.add(keras.Layers.Dense(10))
```

# Sequential Model – Layer objects


- For dense layer, check [https://www.tensorflow.org/api\\_docs/python/tf/keras/layers/Dense](https://www.tensorflow.org/api_docs/python/tf/keras/layers/Dense)

- Arguments:

```
__init__(  
    units,  
    activation=None,  
    use_bias=True,  
    kernel_initializer='glorot_uniform',  
    bias_initializer='zeros',  
    kernel_regularizer=None,  
    bias_regularizer=None,  
    activity_regularizer=None,  
    kernel_constraint=None,  
    bias_constraint=None,  
    **kwargs  
)
```

# Print summary of model

- `model.summary()`
- To produce a digram of the model
- `keras.utils.plot_model(model, to_file='model.png', show_shapes=True, show_layer_names=True)`
- This requires pydot and graphviz installed (pip install pydot graphviz)
- You may need to update graphviz package first (Mac: brew install graphviz/Linux: sudo apt-get install graphviz)





# Compile the model

- Use the method `compile()` of `keras.Model`

- `model.compile(optimizer='adam',  
loss='sparse_categorical_crossentropy', metrics=[ 'accuracy' ])`

- You may set your own optimizer and loss:

- `opt_sgd = keras.optimizers.SGD(lr=0.01, decay=1e-6,  
momentum=0.9)`
  - `loss_ssc = keras.losses.SparseCategoricalCrossentropy()`
  - `model.compile(loss=loss_ssc, optimizer=opt_sgd',  
metrics=[ 'accuracy' ])`

- [https://www.tensorflow.org/api\\_docs/python/tf/keras/optimizers](https://www.tensorflow.org/api_docs/python/tf/keras/optimizers)

- [https://www.tensorflow.org/api\\_docs/python/tf/keras/losses](https://www.tensorflow.org/api_docs/python/tf/keras/losses)

```
compile(  
    optimizer='rmsprop',  
    loss=None,  
    metrics=None,  
    loss_weights=None,  
    sample_weight_mode=None,  
    weighted_metrics=None,  
    target_tensors=None,  
    distribute=None,  
    **kwargs  
)
```

# Fit the training data

- `results = model.fit(train_images, train_labels, batch_size=32, epochs=10, validation_split=0.1)`
- `fit()` method of `keras.Model` returns a History object
- The attribute of `History.history` is a dictionary of training loss, validation loss & training metrics and validation metrics recorded at every epoch
- You can see the keys via `results.history.keys()`
- [https://www.tensorflow.org/api\\_docs/python/tf/keras/Model#fit](https://www.tensorflow.org/api_docs/python/tf/keras/Model#fit)

# Same story with functional API

- Check [2\\_fmnist.py](#)

```
in = keras.layers.Input(shape=(784,), name='input')
h0 = keras.layers.Dense(128, activation='relu', name='hidden0')(in)
h1 = keras.layers.Dense(64, activation='relu', name='hidden1')(h0)
out = keras.layers.Dense(10, activation='softmax', name='output')(h1)

hidden_layers = [h0, h1]
model = keras.Model(inputs=in, outputs=out)
```

- Then you can compile the model with specified optimizer, loss and metrics and fit the model with training data.

# Functional API with branches

- Check [discusion6\\_mnist\\_branches.ipynb](#)
- Functional API can handle models with branches, multiple inputs and outputs
- Use `keras.layers.concatenate([list of layers])` to join branches into a layer

# Load, delete and save models

- `model.save('fmnist_trained.hdf5')` # saves the model
- `del model` # deletes the model
- `model = keras.models.load_model('fmnist_trained.hdf5')`  
# loads the model
- **Save and load model weights**
- `model.save_weights('my_model_weights.hdf5')`  
`model.load_weights('my_model_weights.hdf5')`

# Freeze a layer for fine-tuning

- When you want to update only some layers for new data (fine-tuning the model):
- Pass the trainable argument in the layer object
  - `frozen_layer = keras.Layers.Dense(30, trainable=False)`
- Or set the trainable attribute to False
  - `# hidden_layers[0] is a layer`
  - `hidden_layers[0].trainable = False`

# More on keras.Model.fit()

- **verbose:** 0, 1, or 2. Verbosity mode. 0 = silent, 1 = progress bar, 2 = one line per epoch.
- **validation\_split:** Float between 0 and 1. Fraction of the training data to be used as validation data.
- **shuffle:** Boolean (whether to shuffle the training data before each epoch)

```
fit(  
    x=None,  
    y=None,  
    batch_size=None,  
    epochs=1,  
    verbose=1,  
    callbacks=None,  
    validation_split=0.0,  
    validation_data=None,  
    shuffle=True,  
    class_weight=None,  
    sample_weight=None,  
    initial_epoch=0,  
    steps_per_epoch=None,  
    validation_steps=None,  
    validation_freq=1,  
    max_queue_size=10,  
    workers=1,  
    use_multiprocessing=False,  
    **kwargs  
)
```

# More on keras.Model.fit()

- **callbacks:**

List of [keras.callbacks.Callback](#) instances.

- A **callback** is a set of functions to be applied at given stages of the training procedure. You can use callbacks to get a view on internal states and statistics of the model during training.

```
fit(  
    x=None,  
    y=None,  
    batch_size=None,  
    epochs=1,  
    verbose=1,  
    callbacks=None,  
    validation_split=0.0,  
    validation_data=None,  
    shuffle=True,  
    class_weight=None,  
    sample_weight=None,  
    initial_epoch=0,  
    steps_per_epoch=None,  
    validation_steps=None,  
    validation_freq=1,  
    max_queue_size=10,  
    workers=1,  
    use_multiprocessing=False,  
    **kwargs  
)
```



# callbacks

- Gives us ability to save model at certain stage of training, log loss and other statistics of the model, etc.
- You may use classes provided by `keras.callbacks`
- ```
mcp =  
keras.callbacks.ModelCheckpoint(filepath='/saved/weights.hdf5',  
verbose=1, save_best_only=True)  
model.fit(train_images, train_labels, batch_size=32, epochs=25,  
validation_split=0.1, callbacks=[mcp])
```
- Or write your own callbacks which inherit the methods but with custom stuff.
- See [tf.keras.callbacks](https://www.tensorflow.org/api_docs/python/tf/keras/callbacks/Callback)  
[https://www.tensorflow.org/api\\_docs/python/tf/keras/callbacks/Callback](https://www.tensorflow.org/api_docs/python/tf/keras/callbacks/Callback)